

CSCI 3202: Intro to Artificial Intelligence

Lecture 9: Heuristics

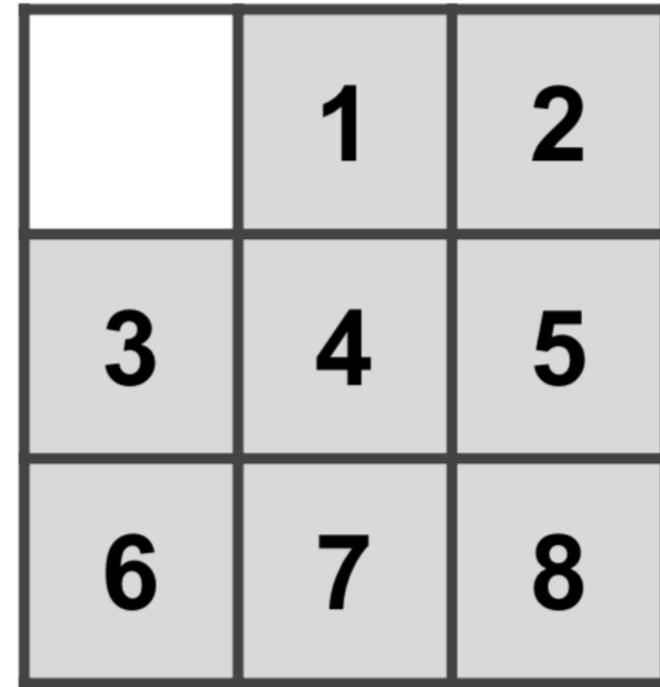
Rachel Cox
Department of
Computer Science



Heuristics

- Using a heuristic can help solve a problem more quickly.
- There is an "art" to deciding on a heuristic function.
- We want $h(n)$ to be admissible. But we need to keep in mind that the lower $h(n)$ is, the more nodes A* expands (making it slower.)

Heuristics

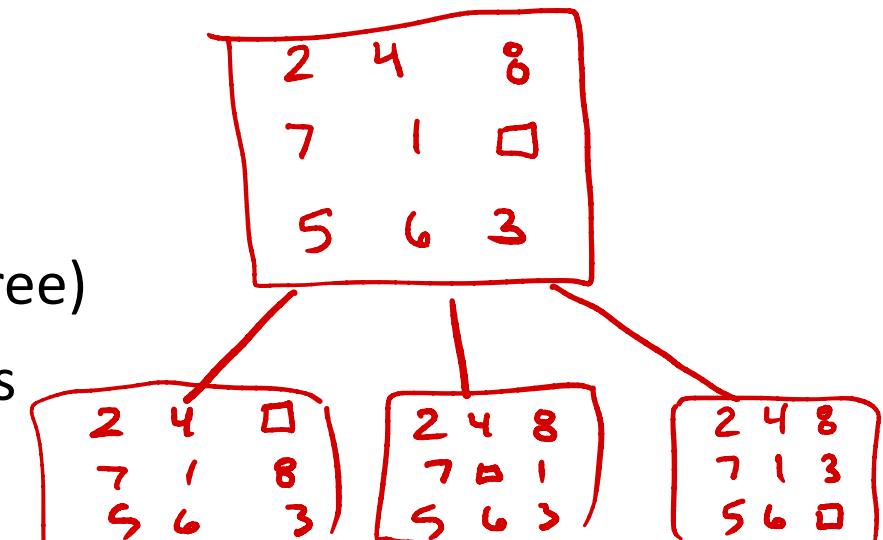


Branching factor $b \approx 3$

Average solution depth = 22

➤ BFS might expand around $3^{22} \approx 3.1 \times 10^{10}$ nodes (tree)

➤ Graph version: $\frac{9!}{2} \approx 180,000$ distinct reachable states



Heuristics

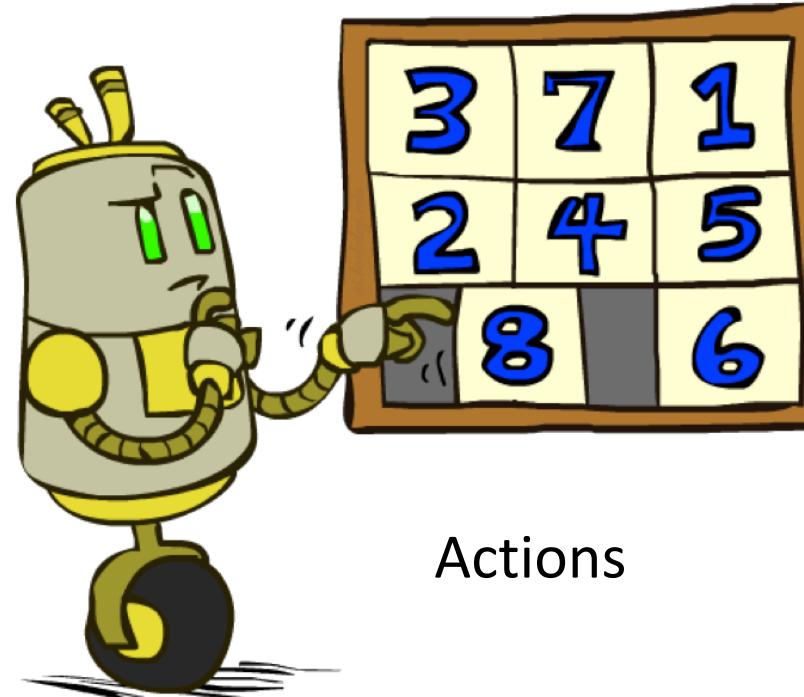
How do we come up with heuristics?

- 1) Generate heuristics from relaxed problems.
- 2) Generate heuristics from sub-problems.
- 3) Learning heuristics from experience.

Heuristics

7	2	4
5		6
8	3	1

Start State



Actions

	1	2
3	4	5
6	7	8

Goal State

- What are the states?
- How many states? 9!
- What are the actions? blank space can slide left , Right , up , down
- How many successors from the start state? 4 successors
- What should the costs be?

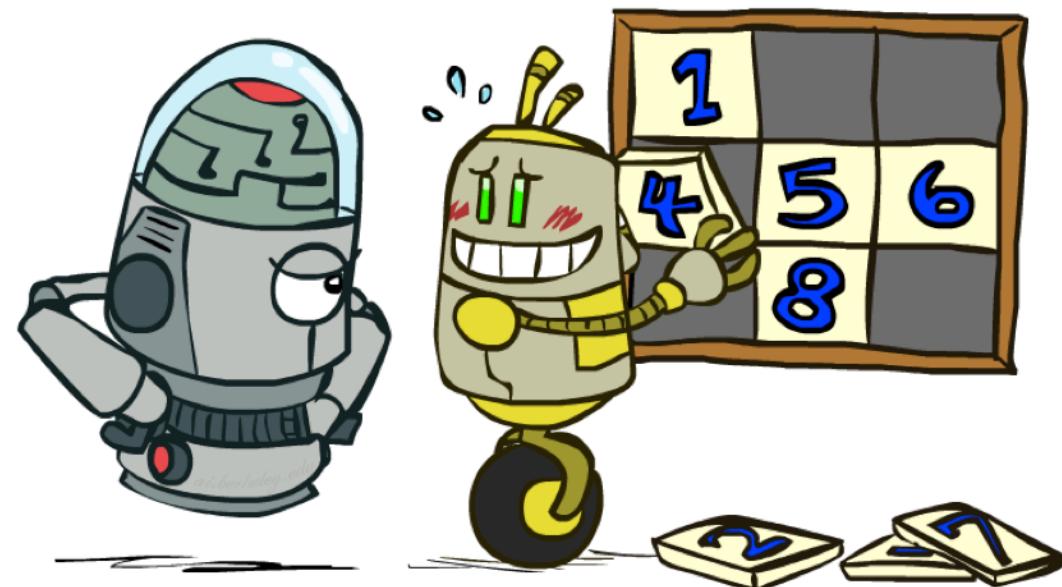
every time we move a tile, we add 1 to the cost

Heuristics

bare minimum
↓

Consider a relaxed problem - allowing the tiles to be pulled off from the board and replaced in correct order

- Heuristic: Number of tiles misplaced
- Why is it admissible?
- $h(\text{start}) = ? \leq 8$
- This is a *relaxed-problem* heuristic



7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8

Goal State



Relaxed problems

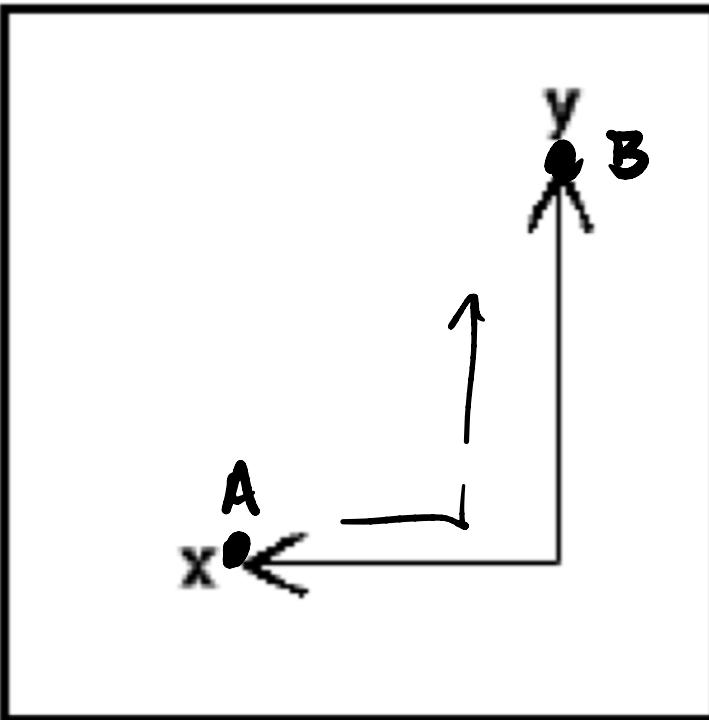
All relaxed problems are doing is adding edges to the state space graph.

⇒ Any optimal solution to the original problem is also a solution to the relaxed problem.

But because of the “short-cuts” of the relaxed problem, there will typically be cheaper solutions than the full problem too.

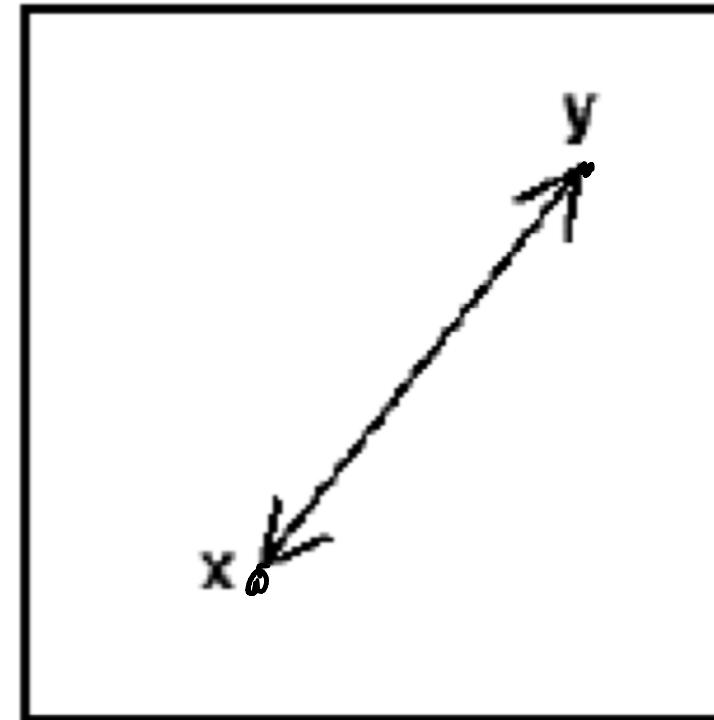
⇒ Optimal solutions of the relaxed problem are admissible heuristics

Heuristics



Manhattan

aka Taxi cab
distance



Euclidean

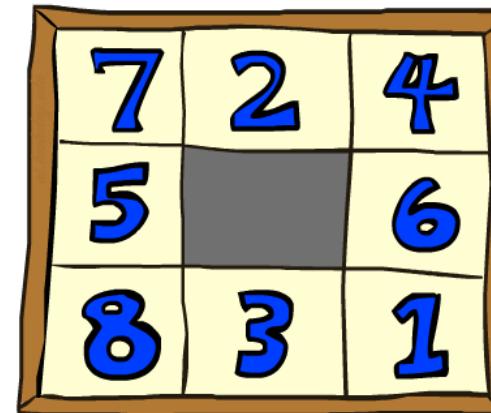
Heuristics

- What if we had an easier 8-puzzle where any tile could slide any direction at any time, ignoring other tiles?

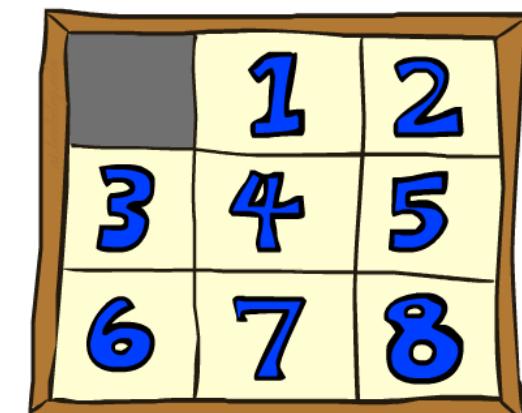
- Total *Manhattan* distance

- Why is it admissible?

$$\begin{aligned} h(\text{start}) &= 3 + 1 + 2 + 2 + 2 + 3 + 3 + 2 \\ &= 18 \end{aligned}$$



Start State



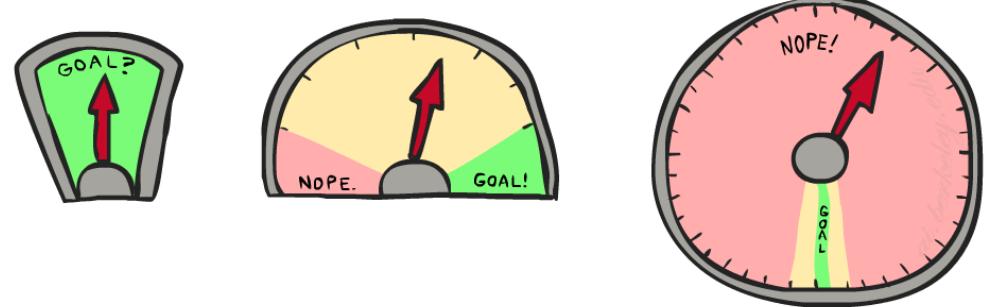
Goal State

Average nodes expanded
when the optimal path has...

TILES	MANHATTAN	4 steps	8 steps	12 steps
13	12	39	25	227

Heuristics

- How about using the *actual cost* as a heuristic?
 - Would it be admissible? $\text{actual cost} \leq \text{actual cost}$
 - Would we save on nodes expanded?
 - What's wrong with it?



- With A*: a trade-off between quality of estimate and work per node
 - As heuristics get closer to the true cost, you will expand fewer nodes but usually do more work per node to compute the heuristic itself

Heuristics

2	4	8
7	1	
5	6	3



	1	2
3	4	5
6	7	8

So we want good heuristics!

- h_1 = number of misplaced tiles
- h_2 = sum of distances of the tiles from their goal positions

Heuristics

Reminder:

Complexity of A*: $O\left((b^\epsilon)^d\right)$

→ The effective branching factor is really b^ϵ

- Suppose A* finds a solution at depth d and generates N nodes to find it.
- b^ϵ is the branching factor that a uniform tree of depth d would have in order to contain $N+1$ nodes:

$$N + 1 = 1 + b^\epsilon + (b^\epsilon)^2 + (b^\epsilon)^3 + \dots + (b^\epsilon)^d$$

Heuristics

Depending on which heuristic we use, h_1 or h_2 , the search cost (nodes generated) and b^ϵ will be different.

Performance comparison

d	Search Cost (nodes)			Effective Branching Factor		
	IDS	A* (h_1)	A* (h_2)	IDS	A* (h_1)	A* (h_2)
2	10	6	6	2.45	1.79	1.79
4	112	13	12	2.87	1.48	1.45
6	680	20	18	2.73	1.34	1.30
8	6384	39	25	2.80	1.33	1.24
10	47127	93	39	2.79	1.38	1.22
12	3644035	227	73	2.78	1.42	1.24
14	-	539	113	-	1.44	1.23
16	-	1301	211	-	1.45	1.25
18	-	3056	363	-	1.46	1.26
20	-	7276	676	-	1.47	1.27

- h_2 (Manhattan distance) dominates h_1 (misplaced tiles)

Heuristics

- A* using h_2 will never expand more nodes than A* using h_1

Every node with $f(n) < C^*$ will be expanded
$$g(n) + h(n) < C^*$$

$\Rightarrow f(n) = g(n) + h(n)$, so every node with $h(n) < C^* - g(n)$ will be expanded

But $h_1(n) \leq h_2(n)$, which means any node expanded by A* using h_2 will be expanded by A* using h_1

So it's best to use a heuristic with higher values

- Makes sense, because those are almost necessarily more accurate:

Admissible \rightarrow Can't overestimate \rightarrow The higher they are, the better they are.

Relaxed problems

Once we have several heuristics, how do we decide which one to use?

- If any dominate, then definitely go with that.
 $\forall n: h_a(n) \geq h_c(n)$
- More generally, just take the most accurate heuristic for whatever node is in question at the moment: $h(n) = \max\{h_1(n), \dots, h_k(n)\}$

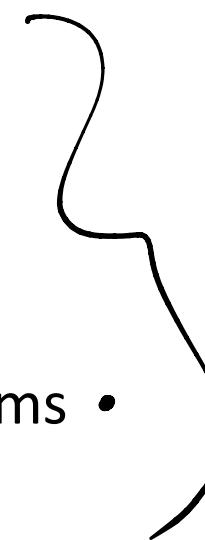


2	4	8
7	1	
5	6	3

Heuristics from sub-problems

Arranging the tiles {1, 2, 3, 4} into the proper slots is a subproblem of the general 8-tile problem.

- The cost of an optimal solution to this subproblem is cheaper than the cost of the optimal solution to the full problem.
- Construct a pattern database:
 - Solve each possible configuration of the subproblem
 - Store the cost of the optimal solution
 - Use this as a heuristic
 - Even better: Do this for multiple subproblems and combine the heuristics



2	4	*
*	1	
*	*	3

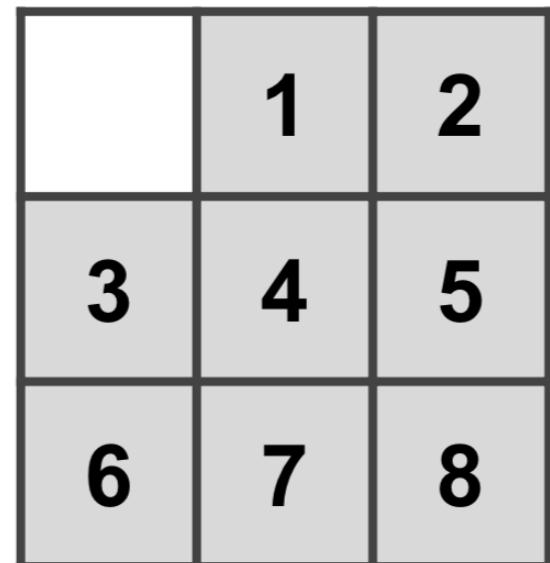
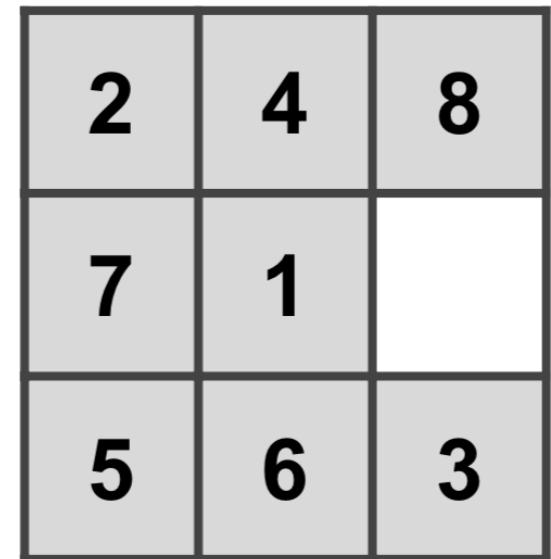
3	4	*
*	*	*

Heuristics from learning

Suppose we solved thousands of 8-tile puzzles

- Then we have a gigantic sample of initial states and of optimal solution paths.

	n_1	n_2	n_3	...
# misplaced tiles ($x_1(n)$)	2	8	5	...
# adjacent tiles that shouldn't be adjacent in goal state ($x_2(n)$)	3	6	4	...
Manhattan distance to goal ($x_3(n)$)	8	14	11	...
Cost	12	24	17	...

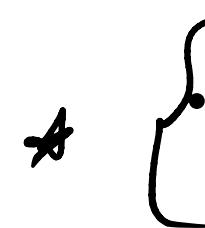


Heuristics from learning

Predict cost from features of the initial states:

$$h(n) = c_1x_1(n) + c_2x_2(n) + c_3x_3(n) + \dots$$

Potential issue:



Not necessarily
admissible/consistent

- Could be, depending on the features and regression constants

	n_1	n_2	n_3	...
# misplaced tiles ($x_1(n)$)	2	8	5	...
# adjacent tiles that shouldn't be adjacent in goal state ($x_2(n)$)	3	6	4	...
Manhattan distance to goal ($x_3(n)$)	8	14	11	...
Cost	12	24	17	...

Next Time

Local Search